

Master Plan Macro Analyzer to Usable State

Created: 2026-06-17 - Owner: Thomas - Branch: `claude/streams-live`

Companion docs: `STATE.md` (live session state) - `DECISIONS.md` (architecture decisions) - `OPEN-QUESTIONS.md` (blockers)

This is the single roadmap to a usable product. `STATE.md` is the per-session handoff; this file is the destination.

1. Definition of "usable"

The operator can sit down and, in one place:

- 1. See every source tiered by conviction (T0->T4), regardless of input type. [DONE] *shipped this session*
- 2. See what each source is actually calling themes, assets, trades with trustworthy per-source accuracy.
- 3. Get a conviction score per source/call that folds in real accuracy (setup-win + alpha), not just trust weight.
- 4. See emerging themes rolled up across all sources (sectors / macro narratives / tickers).
- 5. Get timing signals the core goal: *"when to buy."*
- 6. Run the execution funnel end to end: Positioning -> Concepts -> Identify -> Live -> Journal, on real data.
- 7. Trust durability the listener survives reboots, work is committed, the app is deployable.

We are roughly **5560% there**. Ingestion, extraction, accuracy backtesting, and the source-intelligence UI exist. The gaps are: wiring accuracy -> conviction, themes/timing, and the execution funnel's stubbed data.

2. Current-state snapshot (2026-06-17)

Health: ~565 tests passing, 1 pre-existing failure (`test_watchlist_resolver::mention_summary`). Branch `claude/streams-live` with ~82 uncommitted files (large in-flight effort: insiders, signals, substack, theme-map rewrite + this session's tier/Sources work).

Layer	State
Telegram ingest (5 channels, dedup, albums)	[DONE] Real
Vision extraction (caption-aware, LOCKED prompt)	[DONE] Real 91100% core fields
Full corpus extracted (~6,263 charts)	[DONE] Done
Accuracy backtester (<code>call_accuracy.py</code>)	[DONE] Real <code>setup_win_rate</code> + alpha; 1,494 calls scored
Conviction-tier registry (T0T4 + infra/self)	[DONE] Shipped committed <code>f6ee3a9</code>
Unified Sources view (tier + group clustering, rich drill-down)	[DONE] Shipped this session (uncommitted)
Per-source accuracy badge on <code>TrustedAuthorCard</code>	[PARTIAL] Built but HIDDEN pending trustworthy numbers
Conviction = accuracy-weighted	[TODO] Not wired (still trust-weight only)
Theme map / asset map (live from signals)	[WIP] In-flight (streams-live WIP), 2 tests failing
Timing ("when to buy")	[TODO] Not started the goal
Execution funnel: <code>kpis</code> / <code>heroSignals</code> / <code>watchlist</code>	[TODO] STUB (zero-state) in <code>desk_data.py</code>
Execution funnel: <code>activeTrades</code> / <code>closedTrades</code> / <code>journal</code>	[DONE] Real (empty until first trade)
Durability (launchd listener) + deploy	[TODO] Not started

3. Milestones

Timelines are in **work-sessions** (one focused sitting) plus rough calendar, assuming ~34 sessions/week. Dates are targets, not commitments.

M1 Trust the numbers - *23 sessions - target by 2026-06-24*

Make per-source accuracy trustworthy and visible, then fold it into conviction.

- [] Refine learning/call_accuracy.py: winsorize r_multiple (_MAX_SANE_R=8), min-sample gate (_MIN_SAMPLE=10, meaningful flag), ticker validation (drop junk), market-relative **alpha** (call ret BTC, beta-stripped). *(largely done per STATE verify + lock)*
- [] Re-run backtest_calls() -> source_accuracy(window_days=N); sanity-check the alpha flip (Market Traders +5.0% / Big_Nuts +3.8%).
- [] **Fold accuracy into conviction** in learning/source_themes.py conviction = f(confluence, pattern, TF, persistence, freshness, **setup_win, alpha**).
- [] **Re-enable the accuracy badge** on TrustedAuthorCard (now live in the Sources drill-down) it's currently hidden.
- [] Surface accuracy on the Sources cards (tier grid) so ranking reflects real performance, not just trust weight.
- **Done when:** clicking a T0 source shows its real setup-win-rate + alpha, and conviction scores move with accuracy.

M2 Themes & positioning live - *23 sessions - target by 2026-07-01*

Roll calls up into themes the operator can act on.

- [] Finish the streams-live theme-map/asset-map wiring (real signals -> themeMap[]/assetMap[]); fix the 2 failing test_streams_builders tests (direction/concepts filter).
- [] Conviction-weight the theme rollout (a T0 source's call counts more than a T3's).
- [] Emerging-concepts panel driven by real novelty/velocity.
- [] Decide: does the **Sources** view subsume streams/influence, or do they stay as theme-centric views (IA decision see §5.)
- **Done when:** the operator sees "what's emerging, by conviction" across all sources without reading individual charts.

M3 Timing: "when to buy" - *35 sessions - target by 2026-07-15* - the goal

- [] Define the timing signal: per-ticker, combine (a) conviction-weighted call density, (b) setup freshness/decay, (c) confluence across sources, (d) macro/regime alignment.
- [] Backtest the timing signal against the scored-call outcomes (does "buy when signal fires" beat buy-and-hold / BTC).
- [] Surface timing on the ticker drill-down and as a feed ("act now" candidates).
- **Done when:** the system emits ranked, time-stamped "buy candidates" with a measured edge over baseline.

M4 Execution funnel usable - *34 sessions - target by 2026-07-29*

Wire the stubbed dashboard sections to real data so the funnel is operable.

- [] kpis real cash posture / exposure from trades + regime.
- [] heroSignals[] / watchlist[] real scored setups (currently STUB in desk_data.py).
- [] reasoning{} per-setup explain blob keyed by signalId.
- [] processScorecard real metrics (journal feedback loop already provides the data).
- [] Verify Positioning -> Concepts -> Identify -> Live -> Journal round-trips on real data.
- **Done when:** an idea flows from a source call -> watchlist -> plan -> trade -> review without hitting a stub.

M5 Durability & ship - *23 sessions - target by 2026-08-05*

- [] **launchd plist** for the Telegram listener (survives reboot; replace nohup//tmp/tg-listener.pid).
 - [] **Commit & merge** claude/streams-live to main (currently 82 uncommitted files); land the tier/Sources/group work cleanly.
 - [] Branch/worktree cleanup prune merged claude/* branches and stale worktrees.
 - [] Deployment decision (Render per D-2026-05-08-003) + tactical-gate endpoint live test with Trading-Agent-V1-CODEX.
 - [] Backfill agent_call_log.source_credits column (echo_ties wants it).
 - **Done when:** the app runs unattended, work is on main, and it's deployable.
-

4. Cross-cutting / always-on

- [] Populate Telegram known_senders user-IDs (settings.py) so group posts attribute to individual authors run scripts/telegram_smoke_test.py.
- [] Keep config/sources.json tiers current as new sources are added (CSV -> ingest flow).
- [] Refresh stale market data (prices fetch --watchlist + score run) data has drifted weeks old before.
- [] Keep STATE.md checkpointed at the end of each session; append DECISIONS.md on non-trivial calls.

5. Open decisions (resolve before/while building)

- **IA:** Does **Sources** (tier-grouped roster) fully replace streams + influence, or do those stay as theme/capital-flow views *Leaning: Sources = roster, Streams = themes, retire Influence once Sources covers it.* (M2)
- **Conviction formula weights** confirm the accuracy weighting (how much do setup_win vs alpha count vs the existing confluence/pattern terms). (M1)
- **Timing signal definition** the central modeling choice. (M3)
- **Deployment target** Render still the call Needed for the tactical-gate live test. (M5, from OPEN-QUESTIONS)
- **Group coverage** break out member authors under more groups (e.g. Gem Hunters has none) (cosmetic, anytime)

6. Risks

- **Uncommitted WIP (82 files)** on streams-live is the biggest risk a lost worktree loses weeks. *Mitigation: M5 commit, but consider an interim checkpoint commit sooner.*
 - **Anthropic credits near-\$0** blocks re-extraction only; accuracy/conviction/themes work uses yfinance + DB, so not blocking the critical path.
 - **Listener durability** currently a nohup pid in /tmp; a reboot silently stops forward capture. M5, but low-effort to pull earlier.
-

7. Sequencing summary

```
NOW -- M1 Trust numbers -- M2 Themes -- M3 Timing (GOAL) -- M4 Funnel -- M5 Ship
      (accuracy->conviction) (rollups)   (when to buy)      (un-stub)    (durable+deploy)
      ~by 06-24             ~by 07-01   ~by 07-15          ~by 07-29    ~by 08-05
```

Critical path to "I can act on this": **M1 -> M2 -> M3**. M4 (funnel polish) and M5 (ship) can partially overlap once the intelligence layer is trustworthy.